

FastContentParse - Comprehensive Test Document

§ 1 Scope of Application and General Provisions

This specification applies to all text extraction, document parsing, and layout analysis performed within the FastJava RAG Pipeline framework.

§ 2 Document Parsing Architecture

FastContentParse converts unstructured binary documents into normalized UTF-8 text. It preserves structural metadata such as headings, paragraph boundaries, and page positions. Supported formats include PDF, RTF, Markdown, and plain text.

Section II - Layout Positioning & PDF Extractions

§ 3 Positional Text Stripping

PDF documents store text glyphs with absolute X/Y coordinates on canvas pages. By enabling positional sorting, multi-column articles and sidebars are extracted in natural visual reading order rather than internal PDF stream sequence.

§ 4 Normalization and Cleaning Rules

All Windows-style carriage returns (CRLF) are unified into standard Unix line feeds (LF). Redundant tab stops and vertical feeds are collapsed to maintain clean line margins.

Section III - Vector Indexing & RAG Chunk Pipeline

§ 5 Semantic Chunk Boundary Detection

Extracted text is processed by FastContentChunk to produce contextual passages. Paragraph delimiters and paragraph symbols serve as primary semantic boundaries.

§ 6 Verification and Conformance

This 3-page test document validates multi-page text extraction, page iteration, and RAG pipeline search performance under offline airgap conditions.